

Themelio: a new paradigm for decentralizing the Internet

Yuhao Dong

August 30, 2020

CONTENTS

1	INTRODUCTION	5
1.1	A “blockchain revolution”?	5
1.1.1	The promise of blockchains	5
1.1.2	Where are all the blockchain apps?	6
1.2	What’s wrong with blockchains?	7
1.2.1	Attempts at better blockchains	7
1.2.2	Are blockchains stuck?	8
1.2.3	The crux: endogenous trust	8
1.2.4	Weak endogenous trust in existing block-chains	10
1.3	Towards a new paradigm	11
1.3.1	A minimal blockchain	11
1.3.2	Building a rich ecosystem	12
2	THEMELIO: A MINIMAL BLOCKCHAIN	15
2.1	Design goals	15
2.1.1	Goals	15
2.2	A robust transaction model	16
2.2.1	Coin-based transactions	16
2.2.2	Why coins?	18
2.2.3	Constraint scripting with MelScript	19
2.2.4	Coin-oriented interface	21
2.3	Consensus and trust	22
2.3.1	Oligarchy with a free press	22
2.3.2	Stakeholders: the oligarchy	22
2.3.3	Auditors: the free press	27
2.3.4	Clients: thin yet fully secure	29
2.4	Cryptocurrency and economics	31
2.4.1	Mel: an endogenous stablecoin	31
2.4.2	Better transaction fees	32
3	APPLICATIONS AND PROTOCOLS	35
3.1	Astramel: scalable payments	35
3.2	Conifer and Bitforest: trust-minimizing naming	36
3.3	Token systems	37
3.4	Autonomous applications	38
3.5	Trustless private blockchains	39

INTRODUCTION

1.1 A “BLOCKCHAIN REVOLUTION”?

1.1.1 *The promise of blockchains*

Trust on the Internet is a rare commodity. Participants are often anonymous, and communication is inherently insecure. Everyone is at most a few hundred milliseconds away from potential attackers.

Generally, to trust someone on the internet you either know them in real life or depend on trusted third party intermediaries. However, real-world knowledge is extremely rare on the Internet, while intermediaries (such as certificate authorities, lookup servers, and notaries) are almost always centralized institutions. Unfortunately, central points of trust are often single points of failure. Compromised root CAs lead to catastrophic meltdowns of the basic cryptography of the encrypted Web [12]. Poisoned DNS servers enable defacings of high-profile websites. Hacked update servers can instantly distribute malware to enormous numbers of unsuspecting computers, crippling critical systems [13].

Moreover, centralized control of the “commanding heights” of our modern interconnected society lends a disproportionate amount of power to a small oligarchy of service providers. This enables a wide range of abuse with devastating real-world consequences. For instance, centralized social media platforms insidiously manipulate and censor user communication [15]. As another example, governments build Orwellian surveillance systems like the prototype Chinese “social credit” system [18] by aggregating massive amounts of data from centralized sources.

In the face of these numerous perils of centralization on the Internet, blockchains offer an attractive alternative. Public blockchains like Bitcoin [10] and Ethereum [19] are unforgeable, append-only ledgers accessible to all. They provide secure, transparent, and permanent records of transactions while being completely decentralized. Instead of relying on unaccountable centralized entities, applications such as public key infrastructures, document timestamping services, and electronic money can now use this shared ledger to guarantee security. Blockchains promise a revolution to end reliance on dangerously fragile central trusted entities.

1.1.2 *Where are all the blockchain apps?*

Yet despite all the hype of a blockchain revolution, almost no production systems use public blockchains. Naming even a *single* user-facing application using a public blockchain is astonishingly difficult — except for, of course, cryptocurrency trading apps, blockchain viewers, and other such blockchain-centered software. Why?

Simply put, this is because *public blockchains are not good enough*. Firstly, Nakamoto consensus — the very innovation that gives public blockchains their robust decentralized consensus — saddles users with onerous costs. Take Bitcoin as an example: users must synchronize and store the blockchain’s entire transaction history, which is hundreds of gigabytes and growing. On top of that, users must wait hours for a transaction to become irreversible when the network is a bit slow.

Furthermore, public blockchains can be quite unreliable. Wildly fluctuating cryptocurrency prices create unnecessary currency risk; congestion leads to spikes in transaction fees; and network problems cause long delays. None of these shortcomings are acceptable in modern production systems. Finally, public blockchains by nature require unanimous agreement on the blockchain protocol. Therefore, protocol upgrades are almost always disruptive and contentious — every tiny change can shake the whole blockchain ecosystem. Nobody wants to build applications on a foundation that threatens their products’ basic integrity with every update.

Unsurprisingly, blockchains have failed to truly revolutionize the software industry. Instead, their impact is mostly limited to inspiring a breed of centralized databases with the catchy label of “private blockchains”. Although private blockchains, like those based on Hyperledger Fabric [2], use append-only distributed ledgers similar to those of public blockchains, they’re generally deployed within an environment isolated from public access, such as across a corporate WAN or even inside a single datacenter. Enterprise applications, such as supply-chain tracking or processing business payments, have an especially strong tendency towards using private blockchains.

Unlike their public counterparts, private blockchains promise performance and reliability comparable with traditional databases. For this reason, private blockchains have been adopted by a wide variety of production systems, ranging from the SecureKey identity service [14] to Estonian government systems [6]. “Blockchain for business” is currently almost synonymous to a blockchain within a private, controlled environment.

However, private blockchains by definition give up much of the security, transparency, and decentralization of public blockchains. The majority of proposed deployments, such as those within a single datacenter, almost entirely abandon the promised paradigm shift towards decentralization. Because private blockchains sacrifice decentraliza-

tion in pursuit of maximal usability, they are useless in bringing about a more decentralized Internet.

1.2 WHAT'S WRONG WITH BLOCKCHAINS?

1.2.1 *Attempts at better blockchains*

Many projects attempt to build better blockchains. However, a large number of these propose something categorically different from trustless public blockchains, thus losing sight of the essence of blockchains in pursuit of optimization. Here are two examples.

One fairly obvious idea is to compromise between private and public blockchains to get the best of both worlds. “Consortium” blockchains — blockchains where participation is limited to a small number of trusted partners — are a good example. This includes Quorum [quorum], an Ethereum-based consortium blockchain designed to be deployed in enterprise backend services and Corda [corda], which is designed to provide business-to-business secure transactions. Consortium blockchains may also support public access, though not universal participation in the consensus protocol. Some blockchains, like EOS, even elect the consortium from the wider public.

Limiting the number of consensus participants typically improves performance and reliability simply because the data needs to be replicated dramatically fewer times (this is the the same reason why private blockchains are superior in these aspects). Unfortunately, “hybrid” blockchains retain most of the problems of either public or private blockchains. Those that deliver reliable performance and agile updates end up with the poor security and inadequate decentralization of private blockchains. Others that emphasize security inherit the poor performance and inflexibility of public blockchains. Ultimately, retaining decentralized trust requires wide-scale consensus, while reducing the overhead of consensus necessarily translates into diminished security and neutrality.

A second approach is altogether abandoning the model of a unified, trustless append-only log. This includes “sharded” designs for blockchains, such as Omniledger [kokoris2018omniledger], and non-blockchain ledgers like Ripple [armknecht2015ripple] and Hashgraph [baird2016swirls]. Though forgoing the requirement for a universally replicated log eliminates most scalability challenges, it introduces far greater complexity in both the implementation and application interfaces. This only exacerbates the already daunting difficulty in developing applications on blockchains. In fact, none of these “unconventional” distributed ledgers have achieved much production success. After all, distributed consensus protocols have existed for decades. It is precisely the simple yet expressive abstraction of a linear blockchain

that has fascinated the world, yet unconventional “blockchains” have chosen to throw it away.

1.2.2 *Are blockchains stuck?*

The very existence of multiple “non-blockchain” attempts to fix blockchains hints that public blockchains as conventionally conceived have fundamental problems. Indeed, current blockchains encounter three significant problem areas:

1. **Horizontal scalability:** Global replication and consensus, which are fundamental to the desired blockchain properties, pose severe information-theoretical limits on scaling horizontally. Decentralized apps running on blockchains are already running into this problem, which manifests as expensive and variable fees and poor reliability. In fact, even private and consortium blockchains typically cannot compete with the performance of traditional centralized services.
2. **Governance:** Implementing protocol changes in public blockchains has proven to be fiendishly difficult “political” decisions. They involve much controversy and fail to come to satisfactory results even for simple changes like that of Bitcoin’s block size limit.
3. **Usability:** Blockchains tend to be “leaky” abstractions, requiring developers of blockchain applications to understand low-level details such as block reorganizations and transaction fee auctions. This causes usability problems throughout blockchain applications and often forces end users to deal with the blockchain’s technical problems.

All three problems appear somewhat inherent to the blockchain design and therefore difficult to remedy. Distributed consensus is by nature hard to scale; permissionless blockchains do not have any authority that can “govern” the protocol; blockchains have peculiar characteristics, like Nakamoto consensus, that make usable abstractions hard to build.

Much work has already been done on incrementally attacking these daunting problems, but the issues cannot be wholly eliminated by localized optimization. These problems are symptoms of a more fundamental issue in the very way blockchains are used and designed.

1.2.3 *The crux: endogenous trust*

We cannot build a better blockchain without understanding what exactly makes blockchains unique. Despite their popularity, many of

blockchains' widely-touted features are also found in other systems. For example, *decentralization* is often proclaimed as a key asset of blockchains, but many existing systems have decentralized governance without a central trusted party. This includes DHTs, email, the PGP web of trust, and even the Internet itself. *Transparency* is another commonly-cited advantage of blockchains, yet many non-blockchain protocols like Certificate Transparency and Keybase also use transparency as a key part of their security. Thus, neither decentralization nor transparency constitute the essence of blockchains.

Instead, the crucial feature that distinguishes blockchains from all preexisting protocols is **endogenous trust**. That is, we can trust that a blockchain protocol will behave in a certain way with *minimal assumptions about who runs it*. In blockchains, trust emerges from *within* the protocol, not from preexisting trust in the parties that run the protocol. Crucial to endogenous trust is *cryptoeconomics*, the intersection of game theory and cryptography that enables the formal design of self-incentivizing mechanisms so that given enough participants, we can trust the group's overall behavior without trusting any individual participant.

Endogenous trust is the single most precious property of blockchains that enables secure applications with properties elusive to pre-blockchain systems. For example, someone using a bank must trust the bank no matter what form of communication protocol is used, yet a Bitcoin user does not need to even know which miner ends up processing their transaction, let alone establish some sort of trust with that miner. Instead, the cryptoeconomics of Bitcoin mining internally incentivize miners to cooperate in such a way that makes the Bitcoin protocol as a whole trustworthy.

Many blockchain failures can, in fact, be analyzed as failures in endogenous trust. Contentious governance problems like the Bitcoin block-size controversy of 2017 or the "DAO fork" that split Ethereum into Ethereum and Ethereum Classic arise when external factors incentivize users to override a blockchain's endogenous trust mechanism. Poor internal incentives cause out-of-band coordination to become necessary to prevent game-theoretical issues like "SPV mining" in Bitcoin. Even volatile cryptocurrency prices can be analyzed as a lack of an endogenously trustworthy store of value, forcing users to reach for fiat-pegged stablecoins like Tether that altogether forgo endogenous trust.

Regretfully, the frequency of failure reveals that, in reality, endogenous trust in blockchains is very fragile and often subverted, despite it being the soul of blockchains. Why is that so? An inspection of the state of the art in blockchains reveals the reason.

1.2.4 *Weak endogenous trust in existing blockchains*

Blockchains nowadays can be roughly divided into two categories:

- **Application blockchains** are optimized for one particular application. They often have adequate usability and performance at the expense of adaptability to diverse usages. Examples include Bitcoin Cash (higher throughput bitcoin), Filecoin (incentivized peer-to-peer content distribution), and Zcash (untraceable payments).
- **Platform blockchains** attempt to provide the full set of features needed to implement decentralized applications, typically through Turing-complete “smart contracts”. Ethereum is the archetypal Swiss army knife blockchain; newer examples include EOS and Tezos.

Yet both kinds of blockchains fail at providing strong endogenous trust. This is due to two common problems.

First, *weak cryptoeconomics* often undermine incentive structures intended to secure the blockchain endogenously and force the community to resort to out-of-band coordination to keep the ledger secure. Unpredictable social processes such as Bitcoin’s coordination surrounding SPV mining and EOS’s node elections all result from a lack of built-in incentives nudging participants towards secure behavior. Both application and platform blockchains suffer from poorly designed cryptoeconomic incentives, especially the latter due to their more complex protocols.

The second, and perhaps more important, cause of poor endogenous trust is *application-blockchain friction*. Application-blockchain friction occurs when a blockchain protocol becomes increasingly unsuitable for its main application. When this happens, the blockchain loses its users’ confidence. This forces a contentious out-of-band protocol upgrade to prevent the blockchain from passing into the dustbin of history. The Bitcoin block-size controversy is the most well-known case of loss of trust from application-blockchain friction — unsurprising given Bitcoin’s rigid coupling of its core payment application to its blockchain.

Unfortunately, general-purpose blockchains like Ethereum experience even more challenges to their endogenous trust due to application-blockchain friction. In fact, as we have seen in the previous section, most of the protocol upgrades to Ethereum so far involved fairly minor tweaks to functionality in order to support newly emerging, unanticipated applications.

The prevalence of application-blockchain friction is because both application blockchains (like Bitcoin) and platform blockchains (such as Ethereum) are *are on the wrong protocol layer*. Both are *too close to applications*.

Platform blockchains like Ethereum sit directly underneath applications to allow apps' easy deployment — a new cryptocurrency can be implemented on Ethereum in a few dozen lines of code. Such a direct interface between application and blockchain, however, inevitably results in contention between ever-changing application requirements and ideally immutable blockchain protocols. The situation is analogous to telecommunication networks before the Internet, where a vertically integrated system directly offered relatively high-level functions like voice calling and teletype. Like Ethereum, these complex platforms were extremely costly to upgrade when they were forced to change by the rise of new applications and technological advances.

Learning from previous these mistakes, it is clear that a blockchain with endogenous trust must be built upon a solid cryptoeconomic foundation. This ensures that its core security properties will require no out-of-band social coordination to uphold. More importantly, it must also minimize application-blockchain friction by using a protocol analogous to the Internet Protocol — a minimal, low-level protocol with straightforward semantics. This allows easily upgradable “middleware” protocols to separate the blockchain from the ever-changing needs of applications. Thus, the blockchain can remain an embedded “endogenous trust engine” for decades without changing.

1.3 TOWARDS A NEW PARADIGM

1.3.1 *A minimal blockchain*

This points us towards a new blockchain paradigm — a *minimal blockchain* acting as a bare-bones root-of-trust infrastructure for supporting endogenous trust in decentralized applications. Essentially all other concerns would be subordinated to these two goals.

Current blockchains, as a matter of fact, are very far from the ideal root of trust. Thus, a minimal blockchain design needs to sharply diverge with existing blockchains in the following ways:

- **Minimal governance:** The protocol should be as simple and robust as possible to simply obviate the need for ongoing protocol changes. Deeply embedded infrastructure, such as the Internet Protocol, tend to only be useful if reasonably “timeless”. This is essentially a stronger version of “Szabo’s Law” (“blockchains should not be changed for non-technical reasons”), extended even to technical concerns. Current blockchains, on the other hand, regularly introduce consensus-breaking protocol changes, especially complex platform blockchains like Ethereum.
- **Vertical scalability:** We intentionally avoid pursuing sharding and other horizontal scaling strategies, as they inherently violate logical centralization due to the CAP theorem [brewer2012cap].

Instead, we want a blockchain that effortlessly scales throughput with increasing per-node computational capacity, while supporting fully secure “thin” clients and layer-2 strategies such as state channels at essentially indefinite scale. This is in sharp contrast to present blockchains, where ever-more-subtle shades of eventual consistency are used to support horizontal scaling, protocols treat thin clients as second-class citizens, and most applications are embedded directly in the blockchain state.

- **Cryptoeconomic robustness:** To maximize endogenous trust, we want a blockchain designed with conservative cryptoeconomic assumptions that work without intervention in a wide variety of environments. This is, again, often neglected in current blockchain designs, especially newer projects focused on the putative scalability, governance, or usability issues.
- **Simple abstractions:** Finally, we always choose simple, easy-to-understand abstractions over potentially more powerful but “leakier” ones. For example, we certainly wish to avoid the highly counterintuitive behavior of Nakamoto consensus.

All of these goals attack *systemically infectious* problems in current blockchains when used as a root of trust. No amount of intervening abstraction can protect applications and end users from interventionist governance, sharding-related data inconsistency, cryptoeconomic attacks, and leaky abstractions, all of which can be traced to the application/platform dichotomy and poor endogenous trust. On the other hand, a minimal blockchain that aggressively attacks these systemic problems allows easy encapsulation with well-designed abstractions that hide technical details.

A minimal blockchain, rather than looking like an application or platform, takes inspiration from the technology underpinning most of modern telecommunication: the Internet Protocol (IP). IP gets packets on a best-effort basis from point A to point B, and nothing more. Unreliable datagrams don’t make a developer-friendly interface, but they do provide a firm foundation for ever-changing application protocol stacks. IP is a great illustration of a successful foundational technology. Such protocols are often too simple to support rich applications without intervening protocol stacks, but they are easy to conceptualize, simple to implement, and brutally robust. It is precisely this simplicity that allows it to support the dazzling variety of Internet applications today with practically no changes since the IPv4 specification’s publication in 1981.

1.3.2 *Building a rich ecosystem*

One disadvantage of a minimal blockchain is that it would be difficult to use directly, as it would lack many features developers take for granted,

such as a powerful smart-contract system. A minimal blockchain only makes sense within the context of a richly layered ecosystem, with many abstraction tools available to application developers. Unlike a monolithic platform blockchain, such an ecosystem would be able to rapidly evolve without compromising the blockchain's immutability.

A decentralized-trust ecosystem of applications in this new model can roughly be divided into three layers:

- **Themelio**, a minimal blockchain providing endogenous trust to the entire ecosystem
- **Themelio standard protocols** providing standardization for “middleware” constructs such as state channels, a global naming system, etc that are not hard-coded into the blockchain and may evolve over time
- **Applications** leveraging the Themelio infrastructure to achieve security and decentralization properties impossible without blockchains.

Themelio and its upper-layer protocols concretely instantiate the concept of a minimal blockchain supporting a layered ecosystem. In the process, we investigate and answer three important research questions corresponding to the three divisions above:

- *What is the optimal minimal blockchain?* It seems clear that blockchains should be pushed further away from the application, but what should be the division of functionality between the blockchain and upper-layer protocols? Simply removing “messy” features from existing blockchains will not do, as their entire architecture has not been designed with our paradigm in mind. Exploration of this question has in fact motivated many innovative design choices in Themelio.
- *How to build a “blockchain-minimizing” trustless protocol suite?* Except for a few specific “off-chain” applications such as state channels and cross-chain transfers, research in algorithmic-trust (“trustless”) protocols have largely assumed a blockchain-embedded or similar environment with ubiquitous logical centralization. We need to design network protocols which use our minimal blockchain to provide critical security guarantee, yet “live” largely outside the blockchain to maximize scalability and flexibility.
- *What do applications in such a paradigm look like?* Finally, we explore specific user-facing applications built in this sort of paradigm. We will see that such applications can be exceptionally user-friendly and reliable compared to current blockchain-based applications, while avoiding the governance and security problems inherent with traditional centralized-trust apps.

THEMELIO: A MINIMAL BLOCKCHAIN

In this chapter, the details of Themelio are laid out. Instead of a specific decentralized application or a “Swiss army knife” runtime environment, Themelio provides a *minimal root of trust*. Themelio sits at the bottom of diverse evolving protocol stacks, providing a foundation of endogenous trust and logically centralized consensus — but not much else.

But how are we to build such a blockchain? We combine an elegant, time-tested application interface — a “coin-based” transaction graph like that of Bitcoin — with numerous innovations under the hood. Consensus, based on proof of stake, is immediate, scalable, and highly secure. Cryptocurrency is issued in a completely decentralized manner, yet remains immune to bubbles that destabilize the exchange rate. A simple yet expressive scripting language allows developing advanced decentralized apps without the problems associated with stateful smart contracts. These are some of the many features we use to maximize robustness and performance.

2.1 DESIGN GOALS

Let’s first examine what we want to accomplish, and what we don’t.

2.1.1 Goals

Our overarching goal leads us to design Themelio according to these principles:

1. **Simple abstractions:** Themelio should present simple, “non-leaking” abstractions. Programmers without much experience with blockchains should be able to easily understand the features and behavior of the Themelio blockchain. We try our very best to avoid forcing users to consider subtle edge cases. For example, we must avoid many blockchains’ unintuitive behavior in the presence of network latency (“confirmations”, forks, reorganizations).
2. **Stable protocol:** An initial period of rapid evolution is inevitable. But once mature, Themelio’s internals should change as little as possible. Protocol stability avoids dangerous and messy consensus-breaking updates. Even though many blockchains envision constant protocol evolution, this introduces difficult out-of-band coordination problems — otherwise known

as politics. This can easily lead to de-facto centralization, contentious forks, and subversion by special interests. In Themelio, consensus-breaking changes will be made only in exceptional, non-controversial circumstances, such as to fix critical security vulnerabilities.

3. **Currency stability:** Themelio’s cryptocurrency, the mel, is designed to have very low price volatility. It avoids large price increases, even with spikes in demand. The mel is designed to be a good unit of account and store of value, not a speculative asset exciting to “HODL” (“hold on for dear life”).
4. **High performance:** Themelio’s performance must be much higher than existing public blockchains. This means both high transaction throughput and scalability in the number of fully secure clients. Decentralized apps with debilitatingly poor performance cannot take over the world.
5. **Application neutrality:** Themelio should not attempt to prevent or censor any categories of applications. It does not have an “intended use”.
6. **Robust decentralization:** The ideal public blockchain must simulate a universally trusted intermediary — decentralization is a must. Themelio is designed to decentralize trust across as large a population of stakeholders as possible. No unaccountable third parties, including network operators, should be able to subvert its security guarantees.

2.2 A ROBUST TRANSACTION MODEL

We start with a conceptual exploration of Themelio’s high-level transaction model: the abstractions on the blockchain that applications interact with.

2.2.1 *Coin-based transactions*

Themelio’s basic transaction model belongs to a family usually known as “UTXO-based” or “coin-based” models. This is the oldest family of blockchain models, including first-generation blockchains like Bitcoin and Litecoin. In a coin-based model, the blockchain can be understood as a grow-only directed acyclic graph (DAG) of *transactions*. Every transaction on the blockchain takes as input and spends one or more *unspent transaction outputs (UTXOs)* of previous transactions, which are informally known as *coins*. It then produces as output one or more coins that can be spent as input by subsequent transactions.

Every coin represents a given amount of cryptocurrency, known as its *value*, and it includes an *unlock constraint* that specifies what




Figure 1: Coin-based transactions in Bitcoin

sort of transaction can spend the coin. Each coin can only be spent once. Excepting transactions that “mine” more currency, the sum of the values of all the coins spent by a transaction must equal the sum of the values of all the coins created by it.

Let’s illustrate how coin-based transactions work with a simple example. Assume there are 5 coins identified as B_1, \dots, B_5 , each worth \$1, and each having an unlock constraint specifying “any transaction that spends me must have Bob’s signature”. Informally, we say that Bob owns 5 coins, each one worth \$1. Bob “owning” a coin simply means Bob knowing how to satisfy the coin’s unlock constraint.

Now, assume that Bob wants to send his friend Alice \$2.5. He creates a new transaction spending B_1, B_2, B_3 as input, with two outputs:

- A_1 with value \$2.5 and a constraint requiring Alice’s signature.
- B_6 with value \$0.5 and a constraint requiring Bob’s signature.

and informs Alice about A_1 . Bob now “owns” B_4, B_5, B_6 with a total value of \$2.5, and Alice owns A_1 with a total value of \$2.5, just as we wanted. Note that Bob had to give himself a new coin for the transaction to balance; this new coin is known as a *change output*. Figure 1 from bitcoin.org shows a complex series of interdependent coin-based transactions.

Transactions are batched into an ever-growing series of *blocks*, each one containing transactions settled in a particular time period. Transactions within a block have no defined order — the block that a transaction belongs to is the smallest unit of time on the blockchain. Finally,

blocks are guaranteed to be *consistent*¹, so all users of the blockchain see the same blocks and the same transaction DAG.

2.2.2 Why coins?

In Themelio, we use coin-based transactions with a cryptocurrency that we call the *mel*. We believe that a model of interdependent transactions spending and producing coins, though originally invented only for modeling money transfers, is a very good abstraction on which decentralized-trust applications can be built.

But coin-based models are not popular at all among general-purpose blockchains. Most blockchains attempting to support general decentralized apps use *account and smart-contract* based models. In these models, accounts directly map to sums of money that can be transferred and accounts can have automatically executing code attached. In fact, the only general-purpose blockchain we know of that uses a coin-based model is Qtum. Even there, an “Account Abstraction Layer” simulates Ethereum-like accounts to run smart contracts. Why do we believe coins are the way to go?

First of all, coins allow Themelio to **process transactions quickly**. In an account-based model, like in Ethereum or traditional banking, strict global transaction ordering is necessary. Yet coin-based transactions can be processed in any topological order — we simply need to process the transaction that produces a coin before the transaction that spends it. Transactions within a block can be validated mostly in parallel. This greatly increases performance.

Secondly, a coin-based architecture **simplifies state transitions**. Blockchain protocols can be thought of as state-transition functions, where each transaction takes in the “world” in a certain state (say, Bob having \$5 and Alice \$0) and outputs a different state (Alice and Bob both having \$2.5). To support functionality beyond basic payments, account-based blockchains like Ethereum need arbitrarily mutable global state, accessed by user-programmable “smart contracts”. However, programming decentralized apps with mutable state is notoriously prone to error. Complex state transitions are associated with difficult-to-find bugs and blockchain-level performance problems. In a coin-based blockchain, state is extremely simple: the set of all unspent coins. All transitions simply correspond to individual transactions deleting and adding coins atomically. This leads to clearer logic in decentralized apps and faster performance.

Finally, coin-based transactions are **surprisingly expressive**. A very large class of security-critical problems boil down to establishing a consistent, valid graph of interdependent events. For example, in

¹ Consistency isn’t guaranteed in traditional proof of work blockchains like Bitcoin, but Themelio guarantees immediate, permanent consistency. This is because we use a Byzantine fault-tolerant consensus algorithm, as we will discuss in 2.3

a naming system, a successful name transfer depends on previous events like the previous owner relinquishing control, that owner first registering the name, and so forth. Centralized roots of trust, like notaries, certificate authorities, and banks, almost always serve the role of ensuring consistency of an event graph. In a coin-based blockchain model, the transaction DAG maps extremely well to these event graphs. This means it's easy to write decentralized apps that replicate centralized authorities on Themelio's coin-based model.

However, traditional coin-based architectures exactly like Bitcoin clearly cannot support a wide variety of decentralized apps. Otherwise, why would anybody use other blockchains? Themelio refines the traditional coin-based model with two significant changes: *expressive constraint scripting* and a *coin-oriented application interface*. The former allows programs representing far more than mere “ownership” to constrain coin spending. The latter makes it much easier to write high-performance decentralized apps. Let's now examine these two innovative features of Themelio's transaction model.

2.2.3 Constraint scripting with MelScript

Themelio allows users to write very complex unlock constraints with a powerful scripting language, MelScript. Unlike Bitcoin unlock scripts, MelScript can place conditions on any part of the transaction attempting to spend a coin and enables easy development of a wide variety of decentralized apps. Yet unlike Ethereum's EVM, MelScript is Turing-incomplete and has no access to persistent state, eliminating a large class of “smart contract” bugs.

MelScript is written in a Lisp-like syntax and compiled to a stack-based bytecode to be embedded in transaction outputs. Simple, Bitcoin-like constraints are straightforward. For example, the following is MelScript for a “multisignature” constraint, for coins requiring signatures from both ALICE-KEY and BOB-KEY to be spent:

```
(and (sig-correct? ALICE-KEY)
      (sig-correct? BOB-KEY))
```

We can also access certain facts about the blockchain external to the transaction attempting to spend a coin. For example, the following constraint, which can't be expressed in Bitcoin's simplistic constraint language, gives ownership of a coin to Alice if the total number of transactions exceeds a million before the 10,000th block, and Bob otherwise. It can be used as a simple bet between Alice and Bob on Themelio's future adoption:

```
(if (and (> (get-stat 'transaction-count) 1000000)
        (< (get-stat 'block-height) 10000))
    (sig-correct? ALICE-KEY)
    (sig-correct? BOB-KEY))
```




Figure 2: Example of a Catena log

The most useful constraints in MelScript are not the simple filters demonstrated above, but constraints that constrain constraints. This allows us to embed a wide variety of decentralized, permissionless secure data structures within the transaction graph, which we might call *coin structures*.

This is confusing, so let's illustrate the concept with an example. Catena [17] is an append-only log originally implemented in Bitcoin. The basic idea is simple: a central authority can transparently publish a log of messages by building a transaction chain, each spending the first output coin of the previous transaction. Since coins cannot be spent twice, the authority cannot rewrite, reorder, or delete any log entries after they are published. Figure 2 is an example of a Catena log.

In Bitcoin and other existing coin-based blockchains, Catena logs must be maintained by central authorities. Coins forming the chain must be “owned” by the log publisher, lest someone spend them for other purposes and ruin the log. This prevents the use of Catena in applications without a central publisher. In Themelio, however, we can easily write a MelScript constraint that only allows transactions that grow the Catena chain to spend the coin. Any coin with the following constraint is forced to be the start of a permissionless Catena chain that can never be broken:

```
;; at least 1 output coin
;; first output constrained the same way
;; *SELF* is the coin in which this constraint is embedded
(and (> (output-count) 0)
      (eq? (output-constraint (output-ref 0))
            (output-constraint *SELF*)))
```

Coin structures, of course, are not limited to simple logs. Bitforest [4] builds an entire naming system out of a coin structure that implements an equivocation-proof binary search tree, yet like Catena it must rely on a centralized coin owner when deployed on existing blockchains. Analogous MelScript constraints can be used to implement Bitforest on Themelio as an entirely decentralized and permission-




Figure 3: Application interface of Bitcoin vs Themelio. Squares represent transactions, while circles represent unspent coins.

less naming system, with features comparable to naming systems on “smart contract” blockchains, like the Ethereum Naming System (ENS).

2.2.4 Coin-oriented interface

The second innovation that sets Themelio apart from other blockchains is its deeply coin-oriented application interface. Strange as it may seem, existing coin-based blockchains don’t actually have coins explicitly in the model that applications see. Instead, blockchain users have to download the entire transaction history, building a transaction DAG and working out which transactions spent which coins by themselves. Without full blockchain access, it’s not even possible to securely obtain simple facts like “what are the coins that I own”.

This means that in existing coin-based blockchains, even simple applications like cryptocurrency wallets can’t be secure and scalable at the same time. Either every user downloads the huge and growing transaction history, or a centralized server that does sync the blockchain is trusted to provide users with information ².

In Themelio, though, coins are first-class citizens. Participants synchronize the coin state, not the entire blockchain history. History older than a few weeks is not required to be stored by the protocol. Sparse Merkle trees committing to data about the coin state allow thin clients to securely obtain information about coins without trusting anyone. Apps see the coin state as a secure database they can freely query. Thus, coin-driven applications, ranging from simple wallets to constraint-driven apps like Bitforest, can scale without needing any centralized trust. Figure 3 shows the “worldview” of a Bitcoin node as compared to a Themelio node.

² This is true even with technologies like Bitcoin’s SPV that let clients verify claims that a certain transaction exists. SPV is unable to defend against dishonest nodes that hide coins or claim that spent coins are unspent, rendering basic wallet information untrustworthy.

2.3 CONSENSUS AND TRUST

We now look at how nodes in Themelio come to agreement on the status of the network — decentralized **consensus**, the foundation of any public blockchain’s security. Themelio uses a variation of *bonded proof-of-stake* found in systems such as Tendermint. This is augmented with a novel “auditor” system which further decentralizes trust.

2.3.1 *Oligarchy with a free press*

Participants in Themelio are divided into three categories by their roles:

- **Stakeholder nodes** record transactions into new blocks and confirm them using a Byzantine fault-tolerant consensus algorithm between themselves. They communicate with each other through a broadcast protocol which other nodes in the network never participate in. Anybody can become a stakeholder by “staking” a cryptoasset. Stakeholders correspond to miners or validators in other systems.
- **Auditor nodes** download newly created blocks from the stakeholders and gossip them between themselves while storing a local copy of the coin state. Anybody can join the network as an auditor by simply running a piece of software. Auditors verify new blocks decided by the stakeholders and check that the stakeholders never equivocate on the content of a given block height. Auditors roughly correspond to full nodes in other systems, although they have a more important role in Themelio’s security.
- **Client nodes** are lightweight participants that query the network of auditors to access specific information in the blockchain, yet do not trust any particular auditor.

From this overview we can already see that the trust model of Themelio differs significantly both from that of traditional public blockchains like Bitcoin and from that of typical private blockchains. This is one of its major innovations. Themelio’s trust can be summarized succinctly as an “**oligarchy with a free press**”

2.3.2 *Stakeholders: the oligarchy*

Synkletos: a new approach to proof-of-stake

In Themelio, a Byzantine-resistant fault tolerant algorithm based on HotStuff is used between the stakeholders to establish consensus on the content of the blockchain. The stakeholders form an “oligarchy”:

most users are not stakeholders, yet they get to decide the authoritative state of the network. We call Themelio’s consensus algorithm, together with the cryptoeconomic mechanisms that keep stakeholders honest, *Synkletos*, after the Greek name for the Byzantine Senate.

Synkletos keeps track of a special secondary currency on the blockchain known as the *sym*. Syms are traded freely alongside mels, the main cryptocurrency of Themelio, with a regulated supply of 1 sym per block (1.05 million syms per year). They can be thought of as “shares” in a decentralized corporation in charge of deciding new blocks.

In order to become a stakeholder, one *stakes* at least 1,000 syms, locking them up for a fixed period of time (at least 500,000 blocks, or approximately 6 months) as a performance bond. During that period of time, the stakeholder obtains voting rights in the consensus algorithm in proportion to the amount of syms staked. The central security assumption Themelio uses is that *at least 2/3 of the staked syms are in the hands of honest stakeholders* — a fundamental property of Byzantine-fault-tolerant consensus means we can’t get a better threshold.

Two important questions remain:

- Why do we use proof of stake rather than another consensus mechanism?
- How does *Synkletos* incentivize stakeholders to behave honestly, and what makes *Synkletos*’ approach unique?

Why proof-of-stake?

Synkletos is a variation on *proof of stake* (PoS), a family of blockchain consensus algorithms including Tendermint and Casper. In proof of stake, influence over the consensus process is in proportion to owning an asset, in this case syms. Why did we choose PoS over other consensus algorithms, such as the venerable proof of work (PoW) of Bitcoin, or the proof of authority (PoA) found in consortium and private chains? The Ethereum Proof-of-Stake FAQ [8] give a strong general defense of PoS; we highlight some properties of PoS we consider especially important for Themelio:

- **Higher security margin:** Attacking a PoS blockchain directly requires expending an vast amount of resources to buy up stake. This is equivalent to around 1/3 of the total value of staking coins (a proxy of the economic value of the blockchain system). Thus, as usage increases, PoS security will proportionately strengthen until it becomes practically invulnerable to attacks on the consensus protocol. Proof-of-work blockchains like Bitcoin, however, can be subverted quite cheaply. Attacks reverting a full hour of Bitcoin transactions cost less than \$1,000,000 [11], pocket change

compared to the almost \$100 billion Bitcoin market capitalization. Finally, proof of authority, which is not a decentralized solution, is very fragile to centralized attack vectors such as hacking or government regulation.

- **Immediate finality:** PoS allows easy, secure finality using asynchronous Byzantine fault-tolerant consensus protocols. This means that even if networks are unreliable or malicious, a block that is successfully appended to the blockchain will never be reverted. This eliminates the unpredictable behavior found in “chain-based” consensus protocols like proof of work, such as forks, block reorganizations, and eclipse attacks.
- **Stronger incentive-compatibility:** As we will see later, staked bonds allow us to punish misbehaving stakeholders by deleting their stake. In other blockchains, misbehaving miners only lose potential rewards or reputation. As the Ethereum FAQs put it, “in PoW, we are working directly with the laws of physics. In PoS, we are able to design the protocol in such a way that it has the precise properties that we want - in short, we can optimize the laws of physics in our favor.”

Rewards and slashing

How do we incentivize stakeholders to behave honestly? We use a carrot-and-stick approach commonly found in systems using bonded proof of stake. Honest stakeholders earn *rewards* over time in proportion to the amount of syms they stake, while misbehaving stakeholders can have their entire stake *slashed* given evidence of misbehavior.

Rewards to stakeholders come from two sources: **sym inflation** and **transaction fees**. Stakeholders proposing new blocks earn rewards of 1 newly minted sym per block, just like how Bitcoin miners earn a fixed per-block reward. This implicitly taxes unstaked syms, discouraging holding syms without staking them. Unlike in most other blockchains, inflation is not intended to be the main source of stakeholder income. Since we mint a fixed amount of syms every block, the growth rate in the number of syms as well as the “interest rate” of staking syms approaches zero, making inflation significant only as a short-term bootstrapping subsidy until fees reach a significant level.

Instead, **transaction fees** are used as the main source of stakeholder revenue. Transaction fees, denominated in mels, are imposed on every transaction on the network. Themelio uses a unique mechanism, described in brief in the last section of this whitepaper, to charge a slowly varying uniform fee voted on collectively by the stakeholders. This mechanism, which is crucial to Synkletos’ cryptoeconomic security, allows fees to be a significant and stable source of income for stakeholders, while avoiding well-known game-theoretical attacks

such as fee-stealing associated with conventional auction-based fee markets.

Slashing is the “stick” for punishing cryptographically provable misbehavior. The last step of our Byzantine-fault-tolerant consensus protocol has all stakeholders *commit* to a particular block by signing it cryptographically. Honest stakeholders will always commit a valid block and never “go back” on their collective decision. Thus, we have two *slashing conditions* which leave cryptographic proof that a certain stakeholder is dishonest:

- **Equivocation**, where a stakeholder commits to two different blocks with the same block height
- **Invalid block**, where a stakeholder commits to an invalid block

In either of these cases, anybody can submit cryptographic evidence (two conflicting signatures, or a signature on an invalid block) as a specially-formatted transaction on the blockchain. This **slashing transaction** removes the offending stakeholder, deleting all of the syms associated with the stake. Slashing also reduces the supply of syms and increases the fraction of rewards that other stakeholders receive. This incentivizes large stakeholders to monitor each other and slash misbehaving stakeholders.

Why two currencies?

One of the unique features of Themelio’s proof of stake is its separate staking token, the sym, with features that intentionally discourage use as money. Generally, PoS blockchains use their main “money” coin, like ethers or EOS, as their staking asset. Why not do the same for Themelio? Coins used as stake for consensus are fundamentally equity shares. They are tokens representing fractional ownership of the transaction fees and other “profit” of the system. Unfortunately, *equity shares are a poor form of money*.

First of all, for money we want flexible, demand-responsive monetary policies to reduce value volatility. Otherwise, the currency becomes an unpredictable store of value and a useless unit of account. In the offline world, this is accomplished either by central bank policies for fiat money, or natural supply elasticity for commodity money. For equity shares though, unpredictable share dilution demolishes their fundamental value proposition as fixed slices of profit. Thus, fixed minting schedules, like those of bitcoins and syms, are perfect for equity, but terrible for a new currency.

Furthermore, demand for equity shares in an efficient market is driven largely by speculation on future cash flow, while demand for cash derives from the need for a medium of exchange. We don’t want users of a currency to be forced to speculate on the future transaction fees of a blockchain. Furthermore, increases in currency adoption as a

means of exchange shouldn't drive destabilizing bubbles in currency value.

Themelio therefore uses an independent currency, the *sym*, for the role of equity stake. An independent equity token also turns out to be crucial for establishing backing capital to stabilize the price of mels, as detailed in the Melmint paper.

Achieving high performance

Scalable blockchains with immediate finality need a way to limit the number of consensus participants. This is because Byzantine fault-tolerant consensus algorithms have rapidly increasing overhead with increasing participants. To achieve reasonable scalability and performance, we are forced to limit the number of stakeholders to below a few thousand.

Themelio's way of restricting the number of stakeholders is through the minimum requirement of 1,000 syms staked per validator. Essentially, we limit entry into the oligarchy of stakeholders to only the richest sym holders. Since sym supply follows a fixed schedule, this places a hard limit of a few hundred new validators a year, so that growth in overhead won't outpace growth in computational capacity.

A large minimum stake seems, on the surface, unfair and unappealing. After all, it "disenfranchises" the vast majority of potential stakeholders and institutes a "plutocracy"! Unfortunately, other approaches that superficially sound more decentralized tend to have crippling incentive problems. Ironically, they end up a lot more vulnerable to centralized threats.

For example, a common method of deriving a small amount of participants from a large body of coinholders is *delegated proof of stake* (DPoS). In DPoS, coinholders vote for people with voting power proportional to their coin ownership, and only the few with the most votes become "delegates" and participate in consensus. EOS is a popular blockchain using DPoS.

Yet although DPoS gives a vote to all coinholders, it insulates coinholders from protocol incentives. Coinholders are not responsible for the actions of the delegates they vote for, while misbehaving delegates receive no punishment other than a loss of reputation. Thus, coinholders have no incentive to vote for "good" nodes, delegates have little incentive to behave correctly, and misbehavior is rampant. Unsurprisingly, all the problems of political governance in a representative democracy get imported. Elections involve massive advertising campaigns, vote-buying, and even nationalist agitation[16], while delegates often behave as a centralized cartel, engaging in actions like censoring transactions[7].

Sortition is another approach, used most notably in Algorand [9]. Periodically, a committee of participants is randomly selected from all coinholders — each coinholder has a probability to win this "lottery" in

proportion to the coins that they hold. The committee then participates in a consensus protocol to decide new blocks until the next lottery comes around.

Sortition eliminates most of the politics-like problems of DPoS, allowing protocol incentives like rewards and slashing to work fairly well. Unfortunately, severe problems remain. Randomly selecting participants trustlessly turns out to be a surprisingly hard cryptographic problem — a corrupt lottery can reliably elect malicious committees. Bribery attacks also become much easier, since instead of buying 1/3 of the coins, attackers can simply bribe the current committee, who has only a small fraction of the stake. Complex consensus protocols and advanced, non-quantum-resistant cryptographic techniques can reduce both challenges. But “fancy” mechanisms generally go against Themelio’s philosophy of future-proof simplicity.

A point must be made that *blockchain consensus is not analogous to political governance*. Themelio’s “plutocratic oligarchy” of stakeholders certainly does not make for an effective way of electing a parliament. But for blockchains, it yields highly robust and decentralized security. It disperses control over blockchain consensus to the few hundred people most invested in the health of the network. At the same time, the Synkletos protocol keeps them correctly behaving with massive carrots and sticks. Stakeholders do not decide political questions for the Themelio community; their only job is to run the consensus algorithm correctly.

Thus, we do not believe that Themelio’s “plutocratic” bonded proof of stake is any more vulnerable to centralized threats than PoS blockchains without minimum stake amounts. Even so, Themelio has a system of *auditors* keeping stakeholders in check, ensuring that even a fully corrupted quorum of stakeholders cannot do much damage.

2.3.3 Auditors: the free press

Making failure catastrophic

The “free press” in Themelio consists of *auditors*. Auditors are “full nodes” in usual terminology, replicating and validating the entire blockchain. They form a random *gossip* network among themselves, similar to that used by Bitcoin full nodes. Through this gossip network, information about new blocks is disseminated. Gossip reduces load on the stakeholders and makes it difficult for malicious networks to censor the blockchain — as long as some auditors can connect to the stakeholders and the auditors form a connected graph, new blocks will quickly be visible to every auditor.

The more important role of auditors, though, is to *make consensus failure catastrophic*. This plays a crucial role in keeping the oligarchy of stakeholders honest. Auditors utilize their position as relayers of new blocks to continually monitor for evidence that the stakeholder

consensus is corrupt. For example, invalid blocks or two different blocks at the same height signed by a quorum would be proof that the coordinators are no longer trustworthy. These pieces of evidence, known as *consensus nukes*, undeniably prove that at least $1/3$ of the stakeholders are actively malicious or compromised.

Any auditor that sees a consensus nuke immediately broadcasts it to all auditors it knows in the gossip network. It then permanently activates a “kill switch” and refuses to operate normally. Thus, an attempt at forking or appending invalid transactions to the blockchain would figuratively “nuke” the entire network.

Why consensus nukes?

This objective seems a little strange. Why would we ever want our network to self-destruct?

The obvious answer is that if we no longer have a $2/3$ supermajority of honest stake, the entire system is irrecoverable. More specifically, a well-known result [5] mathematically proves that consensus protocols running in a partially synchronous network model (that is, network delays are unknown but finite) cannot possibly tolerate more than $1/3$ arbitrary faults. So we have to choose between a model where the network stays up, but malicious stakeholders can corrupt the state arbitrarily (rewriting history, giving themselves free money — or shutting down the network), or one where the only thing a corrupted quorum can do is shut down the network. Clearly, the latter is preferable.

More importantly, consensus nuking changes the incentives of potential attackers by making most attacks unprofitable³. Consider a blockchain where consensus-breaking attacks (like Bitcoin’s 51% attack) allow arbitrary state corruption. A malicious actor with the ability to execute such attacks can extract huge profits simply through double-spending. With more complex higher-level applications relying on blockchain data, profit opportunities are even more numerous. Thus, if enough rationally self-serving stakeholders collude, they are greatly incentivized to attack the network and destroy its security guarantees.

If a successful attack can only result in the network stopping all work, only attackers who benefit from destroying the network will participate. Since a successful attacker must stake a vast amount of symets to take over more than $1/3$ of the stake, destroying the network and thus the value of the investment is usually irrational.

Finally, a shutdown when a successful attack occurs forces Themelio users to manually coordinate an emergency “hard fork” out-of-band

³ Consensus nuking can also be seen as a variation on “engineering security through coordination problems”, a concept explored in a blog post by Vitalik Buterin [1]. Attacks by cartels are made impractical because they would require coordinating many users to achieve a cartel-favorable result after the nuke and manual recovery.

to restore the network. This would involve, at the very least, a redistribution of stakes away from the attacking parties and possibly protocol improvements to prevent future attacks. On the other hand, if the blockchain continues to operate even when stakeholders are corrupting the state, nothing forces users to coordinate a hard fork. It's conceivable that the malicious stakeholder cartel can create a climate of pressure for users to go along with the corrupted chain — for example, the state corruption might be forced by legal regulation or presented as way of restoring stolen assets. Consensus nuking ensures that these scenarios are impossible.

2.3.4 *Clients: thin yet fully secure*

Most users of a blockchain, Themelio not excepted, do not have nearly enough resources to process all transactions 24/7. Users that do not synchronize the whole blockchain state, known as thin clients, serve a vital role in any blockchain system. In other blockchains, though, thin clients come with both reduced security and mediocre performance. Bitcoin, for example, has thin clients who must persistently store a growing set of block headers and connect to at least one trusted full node.

In Themelio, thin clients (usually just called clients) are both thinner and safer than thin clients in other systems. Clients only synchronize a small piece of data, less than a kilobyte in size, a few times a year. Yet with this data, they can fully validate a large variety of information they can freely obtain from auditors. Even if a client only connects to bad auditors, it cannot be fooled into accepting invalid data. We accomplish this through two technical innovations: *metastate commitments* in block headers and *epoch-based stake bonds*.

Metastate commitments

Blocks in Themelio, like those in almost all blockchains, have constant-size *headers* that summarize information about that block. Block headers typically cryptographically *commit* to certain pieces of information, such as the transactions within a block, through hash trees and similar mechanisms. Cryptographic commitments allow thin clients to verify claims about the data they commit to without trusting third parties or downloading the entire blockchain.

In traditional coin-based blockchains, block headers commit only to the previous block header and the transactions within the block. This means thin clients can only verify claims that a certain transaction occurred in a block — this is not enough even for basic applications like wallets to be trustless. Account-based blockchains like Ethereum improve on this by committing to the *state*, or all the information

needed to validate new transactions. This allows apps like wallets that rely on querying the state to run trustlessly.

In Themelio, the state is simply the set of all unspent coins. We use a sparse Merkle tree to commit in the block header to a mapping of the coin identifier (hash of the transaction that produced the coin and index of the coin) to coin metadata (value, constraint, etc). This allows thin clients to verify whether or not a coin is spent at a certain block height.

For some simple applications, like verifying a Catena log, this is sufficient. Unfortunately, many applications on coin-based blockchains need to access more than the plain state mapping. For example, wallets would want to know which coins to spend without proofs of payment from all incoming payers.

We therefore commit not just to the state mapping, but also to *metastate*, or metadata about state. Metastate is not strictly necessary for validating blocks, but cryptographic commitments to it in the block header allows more powerful thin clients. This includes information like:

- Number of unspent coins with a certain constraint
- Total number of transactions in the block
- Current block height

Thus, complex coin-oriented applications can trustlessly run on clients that don't need to synchronize anything but the latest block header.

Epoch-based stake bonds

One problem remains: how are clients supposed to get the latest block header? In many blockchains, clients simply synchronize *all* the block headers. Clients would thus use “proof-of-consensus” information (in Bitcoin's case proof of work) embedded in each header to verify the next.

In Themelio, such a strategy would be prohibitively expensive. One of the tradeoffs we made that allows robust proof-of-stake immune to network problems is greatly increased sizes of proofs of consensus. A proof that a block header belongs to the valid blockchain in Themelio requires cryptographic signatures from at least 2/3 of the stakeholders — about 10 KB for a reasonable number of stakeholders. Furthermore, there are just a lot more blocks in Themelio than in Bitcoin. Instead of blocks 10 minutes apart, Themelio produces a block every 30 seconds. This means that in just a year, the block header consensus proofs would amount to more than 10 GB.

To fix this problem, Themelio divides blocks into *epochs* lasting 500,000 blocks, or about half a year. Within each epoch, the list of

shareholders and their respective voting weights stays the same. All stake-related transactions, such as staking and slashing, take effect only at the start of the next epoch. Finally, the last block header of each epoch embeds a *stake document*, which includes the shareholders and voting weights to use when validating blocks in the next epoch.

This means that to validate, say, block header 1,100,000, we simply need the stake document embedded in block 999,999. And to validate that stake document, we just use the stake document in block 499,999. This process repeats until we get to a stake document we already know about.

So clients simply have to catch up on all the new stake documents they missed — 10 KB every 6 months. Afterwards, they can securely validate the latest block header, which then lets them check claims about almost any fact about coins. Such ultra-thin clients allow apps using Themelio to scale on small devices like smartphones, while keeping trust totally decentralized.

2.4 CRYPTOCURRENCY AND ECONOMICS

Finally, we examine the cryptocurrency economy of Themelio, based on a low-volatility currency called the **mel**.

2.4.1 *Mel: an endogenous stablecoin*

Mels are optimized to be the day-to-day transaction currency on Themelio. One can imagine decentralized apps, grocery stores, and peer-to-peer finance to conduct transactions mainly denominated in mel.

The biggest feature of the mel is that it is an **endogenous stablecoin**. This is a completely novel asset class introduced by Themelio, and it means that the mel maintains a stable value without being pegged to any external asset, such as US dollars or gold. Mels maintain their value without non-endogenously-trusted parties present in every other stablecoin system, such as oracles, governance DAOs, and issuers.

The “magic” behind mel through a currency issuance algorithm we published at CryptoEconSys 2020 [3] called Melmint. The details are available in the paper, but the basic objective of Melmint is to peg the value of 1 mel to 1 “DOSC”, a unit that tracks the cost of 24 hours of sequential computation.

By having a stable purchasing power without sacrificing trust, mels would make it much easier to build secure financial assets, transact in cryptocurrency, and protect wealth with Themelio’s endogenous trust.

2.4.2 Better transaction fees

As in Bitcoin and other public blockchains, each transaction in Themelio includes a transaction fee to compensate stakeholders and make flooding attacks costly. Most other blockchains let transaction senders voluntarily decide whatever fee they like; block creators then decide which transactions to include in the limited space within a block. This functions as a pretty fair and efficient first-price auction, since transactions with more fees relative to the burden they pose to the network get higher priority. Unfortunately, auction-based transaction fees paid to whoever included the transaction in a block have several significant problems:

- **Fees are extremely volatile.** When blocks are filled, average fees will vary quite a lot as demand fluctuates. In practice, persistently full blocks is the norm, whether due to demand increase in protocols like Bitcoin where the block size cap is fixed, or due to block producers setting block limits according to demand as in Ethereum. Thus, fees for full blocks are extremely volatile in existing blockchains, often changing as much as 2x within one block interval. This makes for a very poor user experience.
- **Complex client-side fee estimation.** It's far from trivial how much fees to bid in order to get transactions confirmed in a traditional fee market. Wallets need complicated algorithms to estimate the right amount of fee based on looking at unconfirmed transactions — which thin clients can't even securely monitor.
- **Stakeholder incentive problems.** In a proof-of-stake system like Themelio, we want stakeholder income to come primarily from transaction fees. That way, stakes have values in proportion to the value of the system, making attacks harder as usage grows (and damage increases), while giving stakeholders a disincentive to collude to run Themelio into the ground. But a conventional fee market encourages stakeholders to hide transactions from each other — a transaction you include in a block is a transaction fee that I didn't get — leading to all sorts of pathological “fee-stealing” strategies unless stakeholders have a different source of income. This is why Bitcoin and Ethereum rely heavily on inflation, not fees, to reward block producers, but we don't want that in Themelio.

Thus, we abandon the traditional fee auction model in favor of a system inspired by EIP-1559 [eip1559]. Every transaction pays a mel-denominated fee that has two components. A mandatory *base fee* is calculated by multiplying by the *base fee multiplier* the *weight*




Figure 4: Per-block producer rewards (including sym inflation)

of a transaction, a metric that roughly measures its cost. Transaction senders can then add a *tip* above and beyond the base fee.

Every time a new block is created, the stakeholder proposing the block can adjust the base fee multiplier by up to 1% upwards or downwards — the base fee multiplier then reflects the stake-weighted median of the stakeholders' preferences. Base fees, except for an eighth which is burned, are deposited into a special *fee pool* regardless of who included the transaction into the blockchain; the stakeholder creating a block then withdraws a tiny fraction ($1/65536$) of the fee pool. The net effect is that the base fee of a transaction is distributed to all stakeholders regardless of who made the block that contains the transaction.

Tips, on the other hand, are simply paid to the block producer, like fees in traditional blockchains. We expect tips to be a small fraction of total fees, and they give an incentive for block producers to actually include transactions instead of freeloading on a fee pool replenished by other, more honest block producers.

Figure 4 illustrates the flow of funds every time a new block is created.

Why do our changes to the fee market fix its problems? First of all, fee volatility is greatly reduced. When demand fluctuates in the short term, it would be block sizes that fluctuate, not fees. Stakeholders adjust the base fee multiplier to maximize revenue and limit block sizes, but not rigidly at a defined size. One might object that stakeholders can collude to recreate Bitcoin's fee market — by holding down the multiplier to zero, enforcing an unofficial block size limit, and auctioning off block space based on tips. But a rational stakeholder cartel will not do so, as assuming no change in demand, this will simply greatly increase income volatility without increasing expected total income, while in reality volatile fees will probably scare away some users, actually reducing revenue. Incidentally, this is also why we award base fees to stakeholders rather than burning them as in

EIP-1559, since burning base fees will make colluding to create a fee auction highly profitable.

Secondly, clients no longer need complex algorithms to compute fees. The base fee plus a small tip will be sufficient in all cases to get a transaction onto the blockchain as fast as possible. Applications like wallets or payment processors would easily predict the amount of fees needed.

Finally, although we still reward stakeholders with some sym inflation to discourage passive sym-holding, the use of a fee pool ensures that even though stakeholders are rewarded mostly from fees, there's no incentive to steal fees. In fact, one can think of the fee pool as a sort of long-term trust fund, derived from fees, for a stable Bitcoin-like block reward.

Let's now take a look at the sort of protocol and application ecosystem that Themelio supports. Current patterns for blockchain application development are often unsuitable for Themelio due to the lack of features such as smart contracts and sharding. We start from simple to complicated apps, showing how with the help of novel upper-layer abstractions and protocols, decentralization be much more robust and performant.

3.1 ASTRAMEL: SCALABLE PAYMENTS

Obviously, on-chain mel transactions can be used directly to send money. However, as Themelio only processes at most 1000 transactions a second, this does not scale to the extent required for global microtransactions.

Fortunately, a solution already exists: payment channels. Payment channels networks, used in protocols like Lightning Network on Bitcoin or Raiden on Ethereum, route secure cryptocurrency payments through untrusted intermediaries, allowing fully trustless payments without conducting every transaction on the blockchain. PCNs are in fact a primary example of a largely *logically decentralized* protocol with no consistent global state — they work very similarly to the mutual-credit networks underpinning money transfers in the banking system — yet they derive their security ultimately by relying on the logically centralized blockchain as a root of trust.

Existing payment channel constructions for coin-based blockchains, like the Poon-Dryja channels used in Lightning Network, can easily be ported to Themelio. Perhaps more importantly, MelScript allows powerful bidirectional payment channels to be constructed in a straightforward manner without hacks such as the temporary keypairs used in Poon-Dryja channels.

Compared to payment channel networks on blockchains like Bitcoin or Ethereum, a Themelio PCN would provide far better scalability simply due to the faster blockchain — after all, opening and closing payment channels is still limited by blockchain throughput. And compared to the traditional banking network, payment channels give immediate trustless finality, without any authorities that can steal funds or reverse payments. Even on the scalability side, a payment channel network on Themelio will greatly outperform traditional methods in volume and latency, allowing custodian-free microtransactions for applications like paying road tolls.

We have already developed Astrape, a novel PCN construction that achieves both competitive performance and very strong anonymity. A variant of Astrape, Astramel, is under development as a Themelio project to become a standardized scalable asset-transfer scheme for the Themelio blockchain.

In fact, we intend Astramel to be the primary payment protocol for Themelio, rather than raw on-chain transactions. Even for a “basic” functionality like value transfer, we believe that an indefinitely scalable, logically decentralized protocol should be the standard.

3.2 CONIFER AND BITFOREST: TRUST-MINIMIZING NAMING

Using blockchains to implement decentralized and secure naming systems is not a new use case. Blockstack and ENS are examples of blockchain-backed naming systems, which consistently map human-readable identifiers (such as domain names) to security-critical information like public keys without trusting third parties. Legacy naming systems, like DNS and CA-based PKI, are highly centralized and have very fragile security, so even for traditional centralized services like websites, a blockchain-backed, easily deployable naming system can significantly improve security.

It’s not obvious how to embed a naming system in a coin-based blockchain like Themelio. We have developed and published two naming protocols with different design tradeoffs, Conifer [dong2018conifer] and Bitforest [4], for coin-based blockchains. Both encode name bindings as a combination of an on-chain transaction graph and off-chain centralized servers. Notably, both Conifer and Bitforest are *federated* naming systems similar to DNS, where names are issued by diverse authorities, but both use the logically centralized trusted functionality of a blockchain to ensure that such authorities cannot engage in attacks such as impersonation and equivocation.

Themelio’s more powerful coin-based programming paradigm is especially suited for implementing these transaction graph-based naming systems. Notably, by using MelScript instead of simple authority-controlled signature scripts, we can enforce invariants in on-chain transaction graphs without the name-issuing authorities used in Conifer and Bitforest. This allows us to develop variants of Conifer and Bitforest that are entirely decentralized, using technologies such as DHTs to implement the logically decentralized part of the naming system.

Compared to existing blockchain solutions, a Themelio-based naming system would offer higher performance due to the blockchain’s higher throughput, but much more importantly, a paradigm that encourages extensive use of horizontally-scaling off-chain decentralized protocols. Furthermore, Themelio’s powerful thin-client abilities allow deploying blockchain-based naming directly to the smallest of edge de-

vices, rather than relying on trusted gateways like in Blockstack. This removes one of the biggest challenges in deploying naming systems with fully decentralized trust.

3.3 TOKEN SYSTEMS

Another major use case for Themelio is for tokens, like fundraising tokens, cryptokitties, and new cryptocurrencies. Right now, a token is most commonly implemented as a big stateful smart contract on Ethereum, following API standards like ERC-20. This way of implementing a token, unfortunately, is prone to error, often leading to critical security vulnerabilities. Scalability is also a big challenge, as massive stateful smart contracts cannot be easily parallelized even with advanced features like sharding.

In Themelio, on the other hand, custom tokens are ridiculously easy to create. Themelio tokens rely on a special case in its transaction verification logic. Transactions must be balanced by currency — the total values of mels, symets, etc in the input coins spent must be equal to the total values in the outputs — with the exception that an unlimited number of coins, with unconstrained values, can be created with *unlabeled* units. Outputs with unlabeled units will then create coins in the blockchain state with a new unit derived from the unique ID of the transaction.

Thus, a new cryptocurrency token can be created simply by creating any regular transaction while tacking on an additional unlabeled output with the value set to the maximum supply of the new token. This coin's constraint script will then determine the rules of token issuance — no other transactions can create coins with the same unit, since custom tokens are always denominated by the first transaction that created them.

As an example, imagine FooBar wants to create a new token, FooCoin. FooCoin would be sold to the public at a fixed rate of 1000 nanomels per coin. FooBar would broadcast a transaction, say with ID 0xdeadbeef, with an unlabeled output with an inexhaustibly large value (say 2^{64}) and the following constraint:

```
;; first output sends the rest of the FooCoins
;; and repeats this constraint
;; second output sends FooCoin to buyer
;; third output sends mels to FooBar
(and (output-like? 0
  #:currency 0xdeadbeef
  #:constraint (output-constraint *SELF*))
  (output-like? 1
    #:currency 0xdeadbeef)
  (output-like? 2
```

```
#:currency 'nTMEL
#:constraint FOOBAR
#:value (* (output-value (output-ref 0))
           1000))
```

Anybody can then spend this FooCoin-denominated output, diverting some of the FooCoin to himself, leaving the rest with the same constraint, and giving FooBar mels in compensation. FooCoins not encumbered by this constraint freely transact using the same rules as symets and mels do, with no complex smart contracts needed to handle all the cases. Wallet software, payment processors, etc can simply check for coins denominated in “0xdeadbeef” to support FooCoins.

Analogous constraints can be used to implement more complex rules for fungible cryptocurrency tokens. Non-fungible tokens are implemented simply by creating an a new token unit with a coin that has a value of 1. Since 1 cannot be further subdivided, this means that only one coin of that token can ever exist.

3.4 AUTONOMOUS APPLICATIONS

The vast majority of Ethereum-style smart contracts deal with “asset-like” objects like tokens and names and can usually be translated into a simpler and more robust coin-based version on Themelio. One proposed category of decentralized app, though, typically requires a vast amount of state tracking and complex logic and is hard to implement directly on Themelio — ownerless, fully autonomous applications. For example, a smart contract on Ethereum might be an autonomous financial company, negotiating legal contracts, generating packaging its assets into financial products, and pay the bills for a marketing website while hiring people to maintain it, all without human intervention.

Autonomous blockchain entities as described do not really exist except as a concept, and in any case such programs would not be able to scale on Ethereum and other present blockchains due to their poor performance. An autonomous application with security rooted on Themelio would have to implement most of its logic outside the blockchain even if Themelio could provide the requisite performance. For example, the autonomous financial company might run its business logic on traditional cloud services, while accepting payment and paying for its own bills with mels and issuing Themelio-based financial assets. This way, the program would not spam the blockchain with transactions every time its internal state changes. Trustless operation can be achieved by synchronizing state through a trustless private blockchain, as described in the next section.

3.5 TRUSTLESS PRIVATE BLOCKCHAINS

For some applications, nothing short of a new blockchain with its own would do. Traditionally, this would require deploying a new blockchain, private or public, just for use within the application. But this greatly reduces security, as compromising a blockchain formed by consensus between a small number of people is much easier than taking over a public blockchain like Themelio or Ethereum.

One solution is to use a *metachain*, also known as a *virtualchain*, where every time a transaction happens on the smaller blockchain, it is embedded into a corresponding transaction on a public blockchain. Clients of the metachain then scan the entire public blockchain for valid-looking transactions to reach a consensus on the state of the metachain — a very slow process.

Metachains can, of course, be implemented on Themelio, but using Themelio's more expressive features can greatly increase their performance. For example, all transactions claiming to be part of the metachain can be placed in a permissionless Catena log (see 2.2.3). Metachain clients could then avoid scanning through the mass of unrelated Themelio transactions.

If the state transition function of the metachain can be expressed in a short MelScript constraint, Themelio could even enforce the rules of the metachain. This can be combined with using a unique “non-fungible token” inside the Catena log to label the metachain. That way, any thin client can request the transaction with the one and only unspent output denominated in that token, and that transaction is guaranteed to contain the latest state of the metachain.

As described, metachains would be public and permissionless, but similar techniques can be used to secure private blockchains. Data in metachains can be encrypted with a key that only authorized parties know, and the Catena log can have a signature-checking constraint to block unauthorized users from spamming the metachain. Permissioned metachains have a very attractive combination — they inherit the immutability and trustlessness of Themelio, while preventing public access to the contents of the metachain. We believe this is a much better fit for business applications like bid tendering or supply-chain tracking than private or consortium blockchains.

CONCLUSION

Public blockchains, as originally envisioned, herald a fundamental revolution in the way trust works in distributed systems. Unfortunately, they have not seen widespread usage in production systems, outside of a few applications using private blockchains that eschew most of blockchains' distinctiveness. Blockchain development has also run into many serious obstacles, such as scalability and governance.

In this whitepaper, we argued that the main reason for the seeming failure of public blockchains is an incorrect layering paradigm — current blockchains are generally too close to the application layer, forcing complex blockchain implementations on one hand and “leaky”, rigid applications on the other hand. We propose that the correct paradigm for blockchains is that of a minimal root of trust, providing a magic ingredient of endogenous trust to applications that mostly run outside the blockchain.

We described Themelio, a blockchain we developed to support this vision, using many novel technologies and design tradeoffs not seen in current blockchains. We also illustrated the wide range of applications that can be developed using Themelio within a blockchain-minimizing paradigm.

BIBLIOGRAPHY

- [1] Vitalik Buterin. “Engineering Security Through Coordination Problems”. In: (2017). URL: https://vitalik.ca/general/2017/05/08/coordination_problems.html.
- [2] Christian Cachin. “Architecture of the Hyperledger blockchain fabric”. In: *Workshop on Distributed Cryptocurrencies and Consensus Ledgers*. 2016.
- [3] Yuhao Dong and Raouf Boutaba. “Melmint: trustless stable cryptocurrency”. In: *Cryptoeconomic Systems* (2020).
- [4] Yuhao Dong, Woojung Kim, and Raouf Boutaba. “Bitforest: a Portable and Efficient Blockchain-Based Naming System”. In: *2018 14th International Conference on Network and Service Management (CNSM)*. IEEE. 2018, pp. 226–232.
- [5] Cynthia Dwork, Nancy Lynch, and Larry Stockmeyer. “Consensus in the presence of partial synchrony”. In: *Journal of the ACM (JACM)* 35.2 (1988), pp. 288–323.
- [6] E-Estonia. *KSI Blockchain*. URL: <https://e-estonia.com/solutions/security-and-safety/ksi-blockchain/>.
- [7] *EOS’ Blockchain Arbitrator Orders Freeze of 27 Accounts*. 2018. URL: <https://www.coindesk.com/eos-blockchain-arbitrator-orders-freeze-of-27-accounts>.
- [8] *Ethereum Proof of Stake FAQ*. 2019. URL: <https://github.com/ethereum/wiki/wiki/Proof-of-Stake-FAQ> (visited on 03/29/2019).
- [9] Yossi Gilad et al. “Algorand: Scaling byzantine agreements for cryptocurrencies”. In: *Proceedings of the 26th Symposium on Operating Systems Principles*. ACM. 2017, pp. 51–68.
- [10] Satoshi Nakamoto. “Bitcoin: A peer-to-peer electronic cash system”. In: (2008).
- [11] *PoW 51% Attack Cost*. 2019. URL: <https://www.crypto51.app/>.
- [12] J Ronald Prins and Business Unit Cybercrime. “Diginotar certificate authority breach ‘operation black tulip’”. In: *Fox-IT, November* (2011).
- [13] Ronny Richardson and Max North. “Ransomware: Evolution, mitigation and prevention”. In: *International Management Review* 13.1 (2017), pp. 10–21.
- [14] SecureKey. *Building Trusted Identity Networks*. URL: <https://securekey.com/>.

- [15] Cass R Sunstein. *Republic: Divided democracy in the age of social media*. Princeton University Press, 2018.
- [16] *The EOS supernode election: a national struggle worth “hundreds of billions” (in Chinese)*. 2018. URL: <https://zhuanlan.zhihu.com/p/34902188> (visited on 04/01/2019).
- [17] Alin Tomescu and Srinivas Devadas. “Catena: Efficient non-equivocation via bitcoin”. In: *2017 IEEE Symposium on Security and Privacy (SP)*. IEEE. 2017, pp. 393–409.
- [18] Maya Wang. “China’s Chilling ‘Social Credit’ Blacklist”. In: *The Wall Street Journal* 11 (2017).
- [19] Gavin Wood. “Ethereum: A secure decentralised generalised transaction ledger”. In: *Ethereum Project Yellow Paper* (2014).